Bibliography for the Survey Paper: "Consistency Models in Distributed Shared Memory — Theory and Practice"

Hengfeng Wei

April 4, 2015 ∼ April 15, 2015

Contents

1 Bibliography

References

[1] Adve, S. V., and Boehm, H.-J. Memory models: A case for rethinking parallel languages and hardware. *Commun. ACM 53*, 8 (Aug. 2010), 90–101.

[2] Adve, S. V., and Gharachorloo, K. Shared memory consistency models: A tutorial. *Computer 29*, 12 (Dec. 1996), 66–76.

[3] Afek, Y., Attiya, H., Dolev, D., Gafni, E., Merritt, M., and Shavit, N. Atomic snapshots of shared memory. *J. ACM 40*, 4 (Sept. 1993), 873–890.

[4] Ahamad, M., Bazzi, R. A., John, R., Kohli, P., and Neiger, G. The power of processor consistency. In *Proceedings of the Fifth Annual ACM Symposium on Parallel Algorithms and Architectures* (1993), SPAA '93, ACM, pp. 251–260.

[5] Ahamad, M., Neiger, G., Burns, J., Kohli, P., and Hutto, P. Causal memory: definitions, implementation, and programming. *Distrib. Comput. 9*, 1 (Mar. 1995), 37–49.

[6] Aiyer, A., Alvisi, L., and Bazzi, R. A. On the availability of non-strict quorum systems. In *Proceedings of the 19th International Conference on Distributed Computing* (2005), DISC'05, Springer-Verlag, pp. 48–62.

[7] Aiyer, A. S., Alvisi, L., and Bazzi, R. A. Byzantine and multi-writer k-quorums. In *Proceedings of the 20th International Conference on Distributed Computing* (2006), DISC'06, Springer-Verlag, pp. 443–458.

[8] Alglave, J. A formal hierarchy of weak memory models. 178–210.

[9] Anderson, E., Li, X., Shah, M. A., Tucek, J., and Wylie, J. J. What consistency does your key-value store actually provide? In *Proceedings of the Sixth International Conference on Hot Topics in System Dependability* (2010), HotDep'10, USENIX Association, pp. 1–16.

[10] Attiya, H., and Welch, J. *Distributed Computing: Fundamentals, Simulations and Advanced Topics*. John Wiley & Sons, 2004.

[11] Attiya, H., and Welch, J. L. Sequential consistency versus linearizability. *ACM Trans. Comput. Syst. 12*, 2 (May 1994), 91–122.

[12] Bailis, P., Venkataraman, S., Franklin, M. J., Hellerstein, J. M., and Stoica, I. Probabilistically bounded staleness for practical partial quorums. *Proc. VLDB Endow. 5*, 8 (Apr. 2012), 776–787.

[13] Baker, J., Bond, C., Corbett, J. C., Furman, J. J., Khorlin, A., Larson, J., Leon, J., Li, Y., Lloyd, A., and Yushprakh, V. Megastore: Providing scalable, highly available storage for interactive services. In *CIDR 2011, Fifth Biennial Conference on Innovative Data Systems Research, Asilomar, CA, USA, January 9-12, 2011, Online Proceedings* (2011), pp. 223–234.

[14] Boehm, H.-J., and Adve, S. V. Foundations of the c++ concurrency memory model. In *Proceedings of the 2008 ACM SIGPLAN Conference on Programming Language Design and Implementation* (2008), PLDI '08, ACM, pp. 68–78.

[15] Bonnin, D., and Travers, C. α-register. Sept. 2013.

[16] Borowsky, E., and Gafni, E. Immediate atomic snapshots and fast renaming. In *Proceedings of the Twelfth Annual ACM Symposium on Principles of Distributed Computing* (New York, NY, USA, 1993), PODC '93, ACM, pp. 41–51.

[17] Burrows, M. The chubby lock service for loosely-coupled distributed systems. In *Proceedings of the 7th Symposium on Operating Systems Design and Implementation* (2006), OSDI '06, USENIX Association, pp. 335–350.

[18] Calder, B., Wang, J., Ogus, A., Nilakantan, N., Skjolsvold, A., McKelvie, S., Xu, Y., Srivastav, S., Wu, J., Simitci, H., Haridas, J., Uddaraju, C., Khatri, H., Edwards, A., Bedekar, V., Mainali, S., Abbasi, R., Agarwal, A., Haq, M. F. u., Haq, M. I. u., Bhardwaj, D., Dayanand, S., Adusumilli, A., McNett, M., Sankaran, S., Manivannan, K., and Rigas, L. Windows azure storage: A highly available cloud storage service with strong consistency. In *Proceedings of the Twenty-Third ACM Symposium on Operating Systems Principles* (2011), SOSP '11, ACM, pp. 143–157.

[19] Cantin, J. F. The complexity of verifying memory coherence. In *In Proceedings of the fifteenth annual ACM symposium on Parallel algorithms and architectures (SPAA* (2003), ACM Press, pp. 254–255.

[20] Cantin, J. F., Lipasti, M. H., and Smith, J. E. The complexity of verifying memory coherence and consistency. *IEEE Trans. Parallel Distrib. Syst. 16*, 7 (July 2005), 663–671.

[21] Chang, F., Dean, J., Ghemawat, S., Hsieh, W. C., Wallach, D. A., Burrows, M., Chandra, T., Fikes, A., and Gruber, R. E. Bigtable: A distributed storage system for structured data. In *Proceedings of the 7th USENIX Symposium on Operating Systems Design and Implementation - Volume 7* (2006), OSDI '06, USENIX Association, pp. 15–15.

[22] Cooper, B. F., Ramakrishnan, R., Srivastava, U., Silberstein, A., Bohannon, P., Jacobsen, H.-A., Puz, N., Weaver, D., and Yerneni, R. Pnuts: Yahoo!'s hosted data serving platform. *Proc. VLDB Endow. 1*, 2 (Aug. 2008), 1277–1288.

[23] Corbett, J. C., Dean, J., Epstein, M., Fikes, A., Frost, C., Furman, J. J., Ghemawat, S., Gubarev, A., Heiser, C., Hochschild, P., Hsieh, W., Kanthak, S., Kogan, E., Li, H., Lloyd, A., Melnik, S., Mwaura, D., Nagle, D., Quinlan, S., Rao, R., Rolig, L., Saito, Y., Szymaniak, M., Taylor, C., Wang, R., and Woodford, D. Spanner: Google's globally-distributed database. In *Proceedings of the 10th USENIX Conference on Operating Systems Design and Implementation* (2012), OSDI'12, USENIX Association, pp. 251–264.

[24] DeCandia, G., Hastorun, D., Jampani, M., Kakulapati, G., Lakshman, A., Pilchin, A., Sivasubramanian, S., Vosshall, P., and Vogels, W. Dynamo: Amazon's highly available key-value store. In *Proceedings of Twenty-first ACM SIGOPS Symposium on Operating Systems Principles* (2007), SOSP '07, ACM, pp. 205–220.

[25] Dolev, D., and Shavit, N. Bounded concurrent time-stamping. *SIAM J. Comput. 26*, 2 (Apr. 1997), 418–455.

[26] Dutta, P., Guerraoui, R., Levy, R. R., and Chakraborty, A. How fast can a distributed atomic read be? In *Proceedings of the Twenty-third Annual ACM Symposium on Principles of Distributed Computing* (2004), PODC '04, ACM, pp. 236–245.

[27] Fich, F., and Ruppert, E. Hundreds of impossibility results for distributed computing. *Distrib. Comput. 16*, 2-3 (Sept. 2003), 121–163.

[28] Fischer, M. J., Lynch, N. A., and Paterson, M. S. Impossibility of distributed consensus with one faulty process. *J. ACM 32*, 2 (Apr. 1985), 374–382.

[29] Fischer, M. J., and Merritt, M. Appraising two decades of distributed computing theory research. *Distrib. Comput. 16*, 2-3 (Sept. 2003), 239–247.

[30] Garcia-Molina, H., and Barbara, D. How to assign votes in a distributed system. *J. ACM 32*, 4 (Oct. 1985), 841–860.

[31] George, L. *HBase - The Definitive Guide: Random Access to Your Planet-Size Data.* O'Reilly, 2011.

[32] Gibbons, P. B., and Korach, E. Testing shared memories. *SIAM J. Comput. 26*, 4 (Aug. 1997), 1208–1244.

[33] Golab, W., Hurwitz, J., and Li, X. S. On the k-atomicity-verification problem. In *Proceedings of the 2013 IEEE 33rd International Conference on Distributed Computing Systems* (2013), ICDCS '13, IEEE Computer Society, pp. 591–600.

[34] Golab, W., Li, X., and Shah, M. A. Analyzing consistency properties for fun and profit. In *Proceedings of the 30th Annual ACM SIGACT-SIGOPS Symposium on Principles of Distributed Computing* (2011), PODC '11, ACM, pp. 197–206.

[35] Hadzilacos, V., and Toueg, S. A modular approach to fault-tolerant broadcasts and related problems. Tech. rep., 1994.

[36] Haldar, S., and Vidyasankar, K. Constructing 1-writer multireader multivalued atomic variables from regular variables. *J. ACM 42*, 1 (Jan. 1995), 186–203.

[37] Haldar, S., and Vidyasankar, K. On specification of read/write shared variables. *J. ACM 54*, 6 (Dec. 2007).

[38] Herlihy, M. Concurrency and availability as dual properties of replicated atomic data. *J. ACM 37*, 2 (Apr. 1990), 257–278.

[39] Herlihy, M., and Shavit, N. *The Art of Multiprocessor Programming.* Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 2008.

[40] Herlihy, M., and Shavit, N. On the nature of progress. In *Proceedings of the 15th International Conference on Principles of Distributed Systems* (2011), OPODIS'11, Springer-Verlag, pp. 313–328.

[41] Hunt, P., Konar, M., Junqueira, F. P., and Reed, B. Zookeeper: Wait-free coordination for internet-scale systems. In *Proceedings of the 2010 USENIX Conference on USENIX Annual Technical Conference* (2010), USENIXATC'10, USENIX Association.

[42] Jagadeesan, R., and Riely, J. Between linearizability and quiescent consistency - quantitative quiescent consistency. In *Automata, Languages, and Programming - 41st International Colloquium, ICALP 2014, Copenhagen, Denmark, July 8-11, 2014, Proceedings, Part II* (2014), pp. 220–231.

[43] Kawash, J. Y. *Limitations and Capabilities of Weak Memory Consistency Systems.* PhD thesis, Calgary, Alta., Canada, Canada, 2000.

[44] Lakshman, A., and Malik, P. Cassandra: A decentralized structured storage system. *SIGOPS Oper. Syst. Rev. 44*, 2 (Apr. 2010), 35–40.

[45] Lamport, L. A new solution of dijkstra's concurrent programming problem. *Commun. ACM 17*, 8 (Aug. 1974), 453–455.

[46] Lamport, L. Proving the correctness of multiprocess programs. *IEEE Trans. Softw. Eng. 3*, 2 (Mar. 1977), 125–143.

[47] Lamport, L. The mutual exclusion problem: Partii—statement and solutions. *J. ACM 33*, 2 (Apr. 1986), 327–348.

[48] Lee, H., and Welch, J. L. Randomized registers and iterative algorithms. *Distrib. Comput. 17*, 3 (Mar. 2005), 209–221.

[49] Lipton, R., and Sandberg, J. PRAM: A Scalable Shared Memory. Tech. Rep. CS-TR-180-88, Princeton University, Department of Computer Science, 1988.

[50] Lloyd, W., Freedman, M. J., Kaminsky, M., and Andersen, D. G. Don't settle for eventual: Scalable causal consistency for wide-area storage with cops. In *Proceedings of the Twenty-Third ACM Symposium on Operating Systems Principles* (2011), SOSP '11, ACM, pp. 401–416.

[51] Lloyd, W., Freedman, M. J., Kaminsky, M., and Andersen, D. G. Stronger semantics for low-latency geo-replicated storage. In *Proceedings of the 10th USENIX Conference on Networked Systems Design and Implementation* (2013), nsdi'13, USENIX Association, pp. 313–328.

[52] Lynch, N. A hundred impossibility proofs for distributed computing. In *Proceedings of the Eighth Annual ACM Symposium on Principles of Distributed Computing* (1989), PODC '89, ACM, pp. 1–28.

[53] Lynch, N. A., and Shvartsman, A. A. Robust emulation of shared memory using dynamic quorum-acknowledged broadcasts. In *Proceedings of the 27th International Symposium on Fault-Tolerant Computing (FTCS '97)* (Washington, DC, USA, 1997), FTCS '97, IEEE Computer Society, pp. 272–281.

[54] Malkhi, D., Reiter, M., and Wright, R. Probabilistic quorum systems. In *Proceedings of the Sixteenth Annual ACM Symposium on Principles of Distributed Computing* (1997), PODC '97, ACM, pp. 267–273.

[55] Manson, J., Pugh, W., and Adve, S. V. The java memory model. In *Proceedings of the 32Nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages* (2005), POPL '05, ACM, pp. 378–391.

[56] Musial, P., Nicolaou, N., and Shvartsman, A. A. Implementing distributed shared memory for dynamic networks. *Commun. ACM 57*, 6 (June 2014), 88–98.

[57] Petersen, K., Spreitzer, M. J., Terry, D. B., Theimer, M. M., and Demers, A. J. Flexible update propagation for weakly consistent replication. In *Proceedings of the Sixteenth ACM Symposium on Operating Systems Principles* (1997), SOSP '97, ACM, pp. 288–301.

[58] Saito, Y., and Shapiro, M. Optimistic replication. *ACM Comput. Surv. 37*, 1 (Mar. 2005), 42–81.

[59] Shao, C., Welch, J. L., Pierce, E., and Lee, H. Multiwriter consistency conditions for shared memory registers. *SIAM Journal on Computing 40*, 1 (2011), 28–62.

[60] Silberstein, A., Chen, J., Lomax, D., McMillan, B., Mortazavi, M., Narayan, P. P. S., Ramakrishnan, R., and Sears, R. Pnuts in flight: Web-scale data serving at yahoo. *IEEE Internet Computing 16*, 1 (Jan. 2012), 13–23.

[61] Steinke, R. C., and Nutt, G. J. A unified theory of shared memory consistency. *J. ACM 51*, 5 (Sept. 2004), 800–849.

[62] Sumbaly, R., Kreps, J., Gao, L., Feinberg, A., Soman, C., and Shah, S. Serving large-scale batch computed data with project voldemort. In *Proceedings of the 10th USENIX Conference on File and Storage Technologies* (2012), FAST'12, USENIX Association, pp. 18–18.

[63] Taubenfeld, G. On the computational power of shared objects. In *Proceedings of the 13th International Conference on Principles of Distributed Systems* (2009), OPODIS '09, Springer-Verlag, pp. 270–284.

[64] Taubenfeld, G. Weak read/write registers. In *Distributed Computing and Networking, 14th International Conference, ICDCN 2013, Mumbai, India, January 3-6, 2013. Proceedings* (2013), pp. 423–427.

[65] Terry, D. Replicated data consistency explained through baseball. *Commun. ACM 56*, 12 (Dec. 2013), 82–89.

[66] Terry, D. B., Demers, A. J., Petersen, K., Spreitzer, M., Theimer, M., and Welch, B. W. Session guarantees for weakly consistent replicated data. In *Proceedings of the Third International Conference on Parallel and Distributed Information Systems* (1994), PDIS '94, IEEE Computer Society, pp. 140–149.

[67] Terry, D. B., Prabhakaran, V., Kotla, R., Balakrishnan, M., Aguilera, M. K., and Abu-Libdeh, H. Consistency-based service level agreements for cloud storage. In *Proceedings of the Twenty-Fourth ACM Symposium on Operating Systems Principles* (New York, NY, USA, 2013), SOSP '13, ACM, pp. 309–324.

[68] Terry, D. B., Theimer, M. M., Petersen, K., Demers, A. J., Spreitzer, M. J., and Hauser, C. H. Managing update conflicts in bayou, a weakly connected replicated storage system. In *Proceedings of the Fifteenth ACM Symposium on Operating Systems Principles* (1995), SOSP '95, ACM, pp. 172–182.

[69] Torres-Rojas, F. J., Ahamad, M., and Raynal, M. Timed consistency for shared distributed objects. In *Proceedings of the Eighteenth Annual ACM Symposium on Principles of Distributed Computing* (1999), PODC '99, ACM, pp. 163–172.

[70] van Renesse, R., and Schneider, F. B. Chain replication for supporting high throughput and availability. In *Proceedings of the 6th Conference on Symposium on Opearting Systems Design & Implementation - Volume 6* (2004), OSDI'04, USENIX Association, pp. 7–7.

[71] Vogels, W. Eventually consistent. *Commun. ACM 52*, 1 (Jan. 2009), 40–44.

[72] Yu, H., and Vahdat, A. Design and evaluation of a conit-based continuous consistency model for replicated services. *ACM Trans. Comput. Syst. 20*, 3 (Aug. 2002), 239–282.

[73] Zhou, L. Building reliable large-scale distributed systems: When theory meets practice. *SIGACT News 40*, 3 (Sept. 2009), 78–85.